

# Concurrent Quicksort

---

## Implementation

---

The place where this algorithm can be parallelized is the recursive call which happens in two independent different parts of the array [left of pivot and right of pivot].

I parallelize it in two ways:

- creating new processes
- creating new threads

If new processes are created, we need to share the memory [the array] among all processes, so a shared memory region is created.

## NOTE

Having a large number of threads under a single process is bad. If I try to sort 10000 numbers using threads, a lot of threads get created and the cpu is overloaded. So, I have kept a cap of 50000 numbers to call the threaded sorting function

## Code snippets

---

```
int p = partition(arr, l, r);
```

Normal:

```
quickSort(arr, l, p - 1);  
quickSort(arr, p + 1, r);
```

New processes:

```
int p1 = fork();  
int p2;  
  
if(p1 == 0) {  
    quickSort(arr, l, p - 1);  
    exit(0);  
} else {  
    p2 = fork();  
    if(p2 == 0) {  
        quickSort(arr, p + 1, r);  
        exit(0);  
    }  
}
```

```

    } else {
        waitpid(p1, 0, 0);
        waitpid(p2, 0, 0);
    }
}

```

Threaded:

```

a1.arr = a2.arr = a->arr;

a1.l = a->l;
a1.r = p - 1;

a2.l = p + 1;
a2.r = a->r;

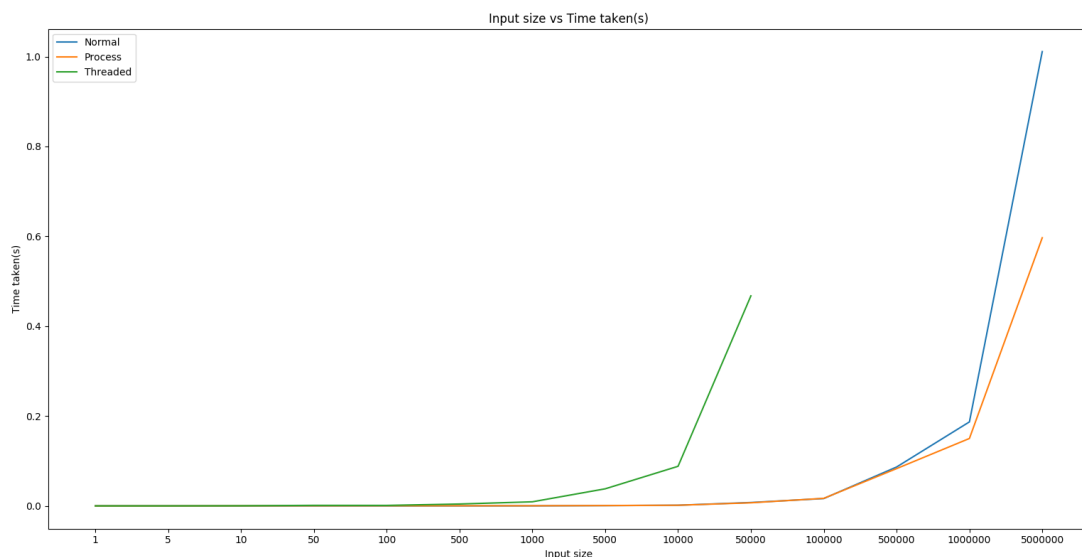
pthread_t t1, t2;
pthread_create(&t1, 0, quickSortThread, &a1);
pthread_create(&t2, 0, quickSortThread, &a2);

pthread_join(t1, 0);
pthread_join(t2, 0);

```

## Results

*A picture speaks a thousand words*



Clearly, it is seen that at first, the parallelizing with process approach is slower than doing the whole thing in a single process, but, as we increase the length of the list, we see that process starts outperforming the single process ones. This starts happening from  $n = 5 \cdot 10^5$

At  $n = 5 \cdot 10^6$ , process takes almost half the time compared to normal